

600.465 — Natural Language Processing

Assignment 2: Probability and Vector Exercises

Prof. J. Eisner — Spring 2016
Due date: Thursday 18 February, 2 pm

Only a bit of programming is required for this assignment. Mostly, you'll solve some pencil-and-paper problems, and then work through a long online visualization. There will be a **short quiz** about the visualization in class on Friday 19 February.

How to hand in your work: Put all notes, documentation, and answers to questions in a README file. The file should be editable so that we can insert comments and give it back to you. For this reason, we strongly prefer a plain ASCII file README, or a L^AT_EX file README.tex (in which case please also submit README.pdf). If you must use a word processor, please save as README.rtf in the portable, non-proprietary RTF format. You should also hand in the code for your findsim program.

Notation: When you are writing the README file, you will need some way of typing mathematical symbols. If your file is just plain ASCII text, please use one of the following three notations and stick to it in your assignment. (If you need some additional notation not described here, just describe it clearly and use it.) Use parentheses as needed to disambiguate division and other operators.

	Text	Picts	L ^A T _E X
$p(x \mid y)$	p(x y)	p(x y)	p(x \mid y)
$\neg x$	NOT x	~x	\neg x
\bar{x} (set complement)	COMPL(x)	\x	\bar{x}
$x \subseteq y$	x SUBSET y	x {= y	x \subseteq y
$x \supseteq y$	x SUPERSET y	x }= y	x \supseteq y
$x \cup y$	x UNION y	x U y	x \cup y
$x \cap y$	x INTERSECT y	x ^ y	x \cap y
$x \geq y$	x GREATEREQ y	x >= y	x \geq y
$x \leq y$	x LESSEQ y	x <= y	x \leq y
\emptyset (empty set)	NULL	0	\emptyset
\mathcal{E} (event space)	E	E	E

1. These short problems will help you get the hang of manipulating probabilities. Let $\mathcal{E} \neq \emptyset$ denote the event space (it's just a set, also known as the outcome space or sample space), and p be a function that assigns a real number in $[0, 1]$ to any subset of \mathcal{E} . This number is called the probability of the subset.

You are told that p satisfies the following two axioms: $p(\mathcal{E}) = 1$. $p(X \cup Y) = p(X) + p(Y)$ provided that $X \cap Y = \emptyset$.¹

¹In fact, probability functions p are also required to satisfy a generalization of this second axiom: if X_1, X_2, X_3, \dots is an infinite sequence of disjoint sets, then $p(\bigcup_{i=1}^{\infty} X_i) = \sum_{i=1}^{\infty} p(X_i)$. But you don't need this for this assignment.

As a matter of notation, remember that the **conditional probability** $p(X | Z) \stackrel{\text{def}}{=} \frac{p(X \cap Z)}{p(Z)}$. For example, singing in the rain is one of my favorite rainy-day activities: so my ratio $p(\text{singing} | \text{rainy}) = \frac{p(\text{singing AND rainy})}{p(\text{rainy})}$ is high. Here the predicate “singing” picks out the set of singing events in \mathcal{E} , “rainy” picks out the set of rainy events, and the conjoined predicate “singing AND rainy” picks out the intersection of these two sets—that is, all events that are both singing AND rainy.

- (a) Prove from the axioms that if $Y \subseteq Z$, then $p(Y) \leq p(Z)$.
You may use any and all set manipulations you like. Remember that $p(A) = 0$ does not imply that $A = \emptyset$ (why not?), and similarly, that $p(B) = p(C)$ does not imply that $B = C$ (even if $B \subseteq C$).
 - (b) Use the above fact to prove that conditional probabilities $p(X | Z)$, just like ordinary probabilities, always fall in the range $[0, 1]$.
 - (c) Prove from the axioms that $p(\emptyset) = 0$.
 - (d) Let \bar{X} denote $\mathcal{E} - X$. Prove from the axioms that $p(X) = 1 - p(\bar{X})$. For example, $p(\text{singing}) = 1 - p(\text{NOT singing})$.
 - (e) Prove from the axioms that $p(\text{singing AND rainy} | \text{rainy}) = p(\text{singing} | \text{rainy})$.
 - (f) Prove from the axioms that $p(X | Y) = 1 - p(\bar{X} | Y)$. For example, $p(\text{singing} | \text{rainy}) = 1 - p(\text{NOT singing} | \text{rainy})$. This is a generalization of [1d](#).
 - (g) Simplify: $(p(X | Y) \cdot p(Y) + p(X | \bar{Y}) \cdot p(\bar{Y})) \cdot p(\bar{Z} | X) / p(\bar{Z})$
 - (h) Under what conditions is it true that $p(\text{singing OR rainy}) = p(\text{singing}) + p(\text{rainy})$?
 - (i) Under what conditions is it true that $p(\text{singing AND rainy}) = p(\text{singing}) \cdot p(\text{rainy})$?
 - (j) Suppose you know that $p(X | Y) = 0$. Prove that $p(X | Y, Z) = 0$.²
 - (k) Suppose you know that $p(W | Y) = 1$. Prove that $p(W | Y, Z) = 1$.
2. All cars are either red or blue. The witness claimed the car that hit the pedestrian was blue. Witnesses are believed to be about 80% reliable in reporting car color (regardless of the actual car color). But only 10% of all cars are blue.

- (a) Write an equation relating the following quantities and perhaps other quantities:

$$\begin{aligned} p(\text{Actual} = \text{blue}) \\ p(\text{Actual} = \text{blue} | \text{Claimed} = \text{blue}) \\ p(\text{Claimed} = \text{blue} | \text{Actual} = \text{blue}) \end{aligned}$$

Reminder: Here, *Claimed* and *Actual* are *random variables*, which means that they are functions over some outcome space. For example, the probability that *Claimed* = blue really means the probability of getting an outcome x such that $\text{Claimed}(x) = \text{blue}$. We are implicitly assuming that the space of outcomes x is something like the set of witnessed car accidents.

²More precisely, $p(X | Y, Z)$ could be either 0 or undefined, namely 0/0. (There do exist advanced ways to redefine conditional probability to avoid this 0/0 problem. Even then, though, one may want a probability measure p to leave some probabilities or conditional probabilities undefined. This turns out to be important for reasons beyond the scope of this course: e.g. http://en.wikipedia.org/wiki/Vitali_set.)

- (b) Match the three probabilities above with the following terms: *prior probability*, *likelihood of the evidence*, *posterior probability*.
- (c) Give the values of all three probabilities. (Hint: Use Bayes' Theorem.) Which probability should the judge care about?
- (d) Let's suppose the numbers 80% and 10% are specific to Baltimore. So in the previous problem, you were implicitly using the following more general version of Bayes' Theorem:

$$p(A | B, Y) = \frac{p(B | A, Y) \cdot p(A | Y)}{p(B | Y)}$$

where Y is $city = \text{Baltimore}$. Just as **1f** generalized **1d**, by adding a “background” condition Y , this version generalizes Bayes' Theorem. Carefully prove it.

- (e) Now prove the more detailed version

$$p(A | B, Y) = \frac{p(B | A, Y) \cdot p(A | Y)}{p(B | A, Y) \cdot p(A | Y) + p(B | \bar{A}, Y) \cdot p(\bar{A} | Y)}$$

which gives a practical way of finding the denominator in the question **2d**.

- (f) Write out the equation given in question **2e** with A , B , and Y replaced by specific propositions from the red-and-blue car problem. For example, Y is “ $city = \text{Baltimore}$ ” (or just “Baltimore” for short). Now replace the probabilities with actual numbers from the problem, such as 0.8.

Yeah, it's a mickeymouse problem, but I promise that writing out a real case of this important formula won't kill you, and may even be good for you (like, on an exam).

3. Beavers can make three cries, which they use to communicate. **bwa** and **bwee** usually mean something like “come” and “go” respectively, and are used during dam maintenance. **kiki** means “watch out!” The following **conditional probability table** shows the probability of the various cries in different situations.

$p(cry situation)$	Predator!	Timber!	I need help!
bwa	0	0.1	0.8
bwee	0	0.6	0.1
kiki	1.0	0.3	0.1

- (a) Notice that each column of the above table sums to 1. Write an equation stating this, in the form $\sum_{variable} p(\dots) = 1$.
- (b) A certain colony of beavers has already cut down all the trees around their dam. As there are no more to chew, $p(timber) = 0$. Getting rid of the trees has also reduced $p(predator)$ to 0.2. These facts are shown in the following **joint probability table**. Fill in the rest of the table, using the previous table and the laws of probability. (Note that the meaning of each table is given in its top left cell.)

$p(\text{cry}, \text{situation})$	Predator!	Timber!	I need help!	TOTAL
bwa				
bwee				
kiki				
TOTAL	0.2	0		

(c) A beaver in this colony cries **kiki**. Given this cry, other beavers try to figure out the probability that there is a predator.

- This probability is written as: $p(\text{_____})$
- It can be rewritten without the $|$ symbol as: _____
- Using the above tables, its value is: _____
- Alternatively, Bayes' Theorem allows you to express this probability as:

$$\frac{p(\text{_____}) \cdot p(\text{_____})}{p(\text{_____}) \cdot p(\text{_____}) + p(\text{_____}) \cdot p(\text{_____}) + p(\text{_____}) \cdot p(\text{_____})}$$

- Using the above tables, the value of this is:

$$\frac{\text{_____} \cdot \text{_____}}{\text{_____} \cdot \text{_____} + \text{_____} \cdot \text{_____} + \text{_____} \cdot \text{_____}}$$

This should give the same result as in part iii., and it should be clear that they are really the same computation—by constructing table (b) and doing part iii., you were *implicitly* using Bayes' Theorem. (I told you it was a trivial theorem!)

4. A **language model** is a probability function p that assigns probabilities to word sequences such as $\vec{w} = (\text{i}, \text{love}, \text{new}, \text{york})$.

Suppose $\vec{w} = w_1 w_2 \cdots w_n$ (a sequence of n words). A **trigram language model** defines

$$p(\vec{w}) \stackrel{\text{def}}{=} \prod_{i=1}^{n+1} p(w_i | w_{i-2}, w_{i-1}) \quad (1)$$

where by convention we define $w_{n+1} = \text{EOS}$ (“end of sequence”) and $w_0 = w_{-1} = \text{BOS}$ (“beginning of sequence”).

For example, the model says that a speaker who says “i love new york” has generated the sequence in left-to-right order, by rolling 5 dice that happened to land on **i**, **love**, **new**, **york**, EOS. Each roll uses a special die selected by the two previously rolled characters (or BOS).

- (a) Expand the above definition of $p(\vec{w})$ using naive estimates of the parameters, such as

$$p(w_4 | w_2, w_3) \stackrel{\text{def}}{=} \frac{c(w_2 w_3 w_4)}{c(w_2 w_3)} \quad (2)$$

where $c(w_2 w_3 w_4)$ denotes the count of times the trigram $w_2 w_3 w_4$ was observed in a training corpus. What property of the corpus is counted by each of the following: $c(\text{BOS BOS})$, $c(\text{BOS BOS i})$, $c(\text{new york EOS})$?

Remark: Naive parameter estimates of this sort are called **maximum-likelihood estimates** (MLE). They have the advantage that they maximize the probability (equivalently, minimize the perplexity) of the training data. But they will generally perform badly on test data, unless the training data were so abundant as to include all possible trigrams many times. This is why we must smooth these estimates in practice.

- (b) In the data you will use in the *next* assignment, a sequence may be very long. While the sequence is still preceded by BOS and terminated with EOS, each linguistic sentence within the sequence is delimited by `<s>` at the start and `</s>` at the end. For example, the following single sequence consists of 3 linguistic sentences:

`<s> do you think so </s> <s> yes </s> <s> at least i thought so </s>`

Given English training data, the probability of

`<s> do you think the </s>`

should be extremely low under any good language model. Why? In the case of the trigram model, which parameter or parameters are responsible for making this probability low?

- (c) You turn on the radio as it is broadcasting an interview. Assuming a trigram model, match up expressions (A), (B), (C) with descriptions (1), (2), (3):

The expression

(A) $p(\text{Do}) \cdot p(\text{you} \mid \text{Do}) \cdot p(\text{think} \mid \text{Do}, \text{you})$

(B) $p(\text{Do} \mid \text{<s>}) \cdot p(\text{you} \mid \text{<s>, Do}) \cdot p(\text{think} \mid \text{Do}, \text{you}) \cdot p(\text{</s>} \mid \text{you}, \text{think})$

(C) $p(\text{Do} \mid \text{<s>}) \cdot p(\text{you} \mid \text{<s>, Do}) \cdot p(\text{think} \mid \text{Do}, \text{you})$

represents the probability that

- (1) the first complete sentence you hear is `Do you think` (as in, "D'ya think?")
- (2) the first 3 words you hear are `Do you think`
- (3) the first complete sentence you hear starts with `Do you think`

Explain your answers briefly. Which quantity is $p(\vec{w})$? *Remark:* The distinctions matter because "Do" is more probable at the start of an English sentence than in the middle, and because (3) describes a larger event set than (1) does.

- (d) **Extra credit:** One could also define a kind of reversed trigram language model p_{reversed} that instead assumed the words were generated in reverse order ("from right to left"):

$$p_{\text{reversed}}(\vec{w}) \stackrel{\text{def}}{=} \prod_{i=0}^n p(w_i \mid w_{i+1}, w_{i+2}) \quad (3)$$

where by convention we define $w_0 = \text{BOS}$ and $w_{n+1} = w_{n+2} = \text{EOS}$. In this case, the generating process stops when we hit BOS, instead of EOS.

By manipulating the notation, show that the two models are identical (i.e., $p(\vec{w}) = p_{\text{reversed}}(\vec{w})$ for any \vec{w}) provided that both models use MLE parameters estimated from the same training data (see problem 4a).

Hint: Try writing out the probability of "i love new york" under both models. To argue that the resulting probabilities are equal, you will have to observe that certain counts are equal.

5. Under an bigram language model, knowing an early part of the sentence (“*Horses like ...* ”) doesn’t tell you too much about the end of the sentence. But that overlooks a useful property of real language data. A sentence that starts out “*Horses like ...* ” is probably about horses, so it will probably continue to have a lot of horse-related words throughout, even at the end.

Suppose your corpus seems to cover k different topics—horse-racing, Star Trek, knitting, Wall Street, etc. For the most part, each sentence seems to stick to a single topic. **You want to build a better bigram model that also captures the fact the choice of topic persists throughout each sentence.**

Write a formula for $p(w_1 w_2 w_3 w_4)$ under this better bigram model. (*Hints:* Latent variable, chain rule, backing off using a conditional independence assumption. Make w_4 depend not only on w_3 but also on the topic a .)

6. So far in this assignment, we have been using conditional probabilities like $p(\text{Actual} = \text{blue})$ or $p(\text{singing} \mid \text{rainy})$ or $p(w_4 \mid w_2, w_3)$. But where do those probabilities come from?

One option would be to estimate conditional event probabilities naively from data, using simple count ratios as in formula (2). But those unsmoothed estimates will not be reliable when the numerator or denominator counts are small.

An alternative is to use log-linear models. Frank Ferraro and I built an online toy to let you play with such models. It should give you almost *physical* intuitions about how these models behave.

You’ll need to understand log-linear models in subsequent lectures and on the exams. They are also a pretty good introduction to machine learning.

You can find the toy at <http://cs.jhu.edu/~jason/465/hw-prob/loglin/>. Work through its 18 lessons (and read the accompanying handout). This should be fun, although it may take a few hours.

To save you time, there is nothing to hand in for this part of the homework. **Instead, there will be a brief in-class quiz** to check that you understood what’s going on. Your score on the quiz will be your score for this homework question.

The lessons are peppered with guidance in the form of questions. Just think through these questions as you go along, and you should do fine on the quiz. We suggest you **work through the lessons with a friend or two**, so that you can discuss them together.

7. Now, how might you use log-linear models in linguistics? Rather than answering in your README, do one of the following:

- Think over your own interests in linguistics, NLP, or a related field. Post a new note on Piazza with a subject line “Log-linear model for [my idea].”

In your note, give an example of a conditional distribution $p(y \mid x)$ that would be interesting to model. Your goal here might be to predict y from x , or to understand the properties of x that are predictive of y . Make sure to be clear about your problem and the model you’d build. What do x and y represent?

- **Alternatively**, respond to someone else’s Piazza post that you found interesting. Suggest some new features $f_k(x, y)$ that would be useful for their problem. Or suggest changes to the problem setup (e.g., the choice of x and y) or to other people’s features.

8. The next assignment will involve various kinds of smoothed language models. Often, smoothing involves treating similar events in similar contexts as having similar probabilities. (Backoff is one example.)

One strategy will be to treat similar *words* as having similar probabilities. But what does “similar words” mean? More concretely, how will we *measure* whether two words are similar?

Log into the ugrad machines, and look in the directory `/usr/local/data/cs465/hw-lm/lexicons/`. Remember that a *lexicon* lists useful properties of the words of a language. Each of the `words-*.txt` files³ is a lexicon that lists over 70,000 words of English, with a representation of each word as a d -dimensional vector. In other words, it *embeds* these words into the vector space \mathbb{R}^d .

We can now define “similar words” as words with similar vectors. A word may have many syntactic and semantic properties—some words are transitive verbs, some words are plural, some words refer to animals. These properties can be considered by a log-linear model, and words with similar vectors tend to have similar properties. The larger d is, the more room the vector has to encode interesting properties.

The vectors in these particular files were produced automatically by Google’s `word2vec` program.⁴ Their individual dimensions are hard to interpret as natural properties. It would certainly be nice if dimension 2 (for example) represented the “animal” property, so that a word that referred to an animal would be one whose vector $\vec{v} = (v_1, v_2, \dots, v_d)$ had strongly positive v_2 . In practice, however, `word2vec` chooses an arbitrary basis for the vector space. So the “animal” direction—to the extent that there is one—is actually represented by some arbitrary vector \vec{u} . The animal words tend to be those whose vectors \vec{v} have strongly positive $\vec{v} \cdot \vec{u}$.

(Geometrically, this dot product measures how far \vec{v} extends in direction \vec{u} (it projects \vec{v} onto the \vec{u} vector). Algebraically, it computes a certain linear combination of v_1, v_2, \dots, v_n . As a special case, if \vec{u} were $(0, 1, 0, 0, \dots)$, then $\vec{v} \cdot \vec{u}$ would in fact return v_2 . But there’s no reason to expect that \vec{u} would be that simple.)

You’ll be using these lexicons in the next assignment. For this assignment, you’ll just warm up by writing a short program to get a feel for vector embeddings.

- (a) Your program `findsim` should print the 5 words most similar to `seattle`, other than `seattle` itself, according to the 50-dimensional embeddings, if you run it as

```
findsim words-50.txt seattle
```

They should be printed in decreasing order of similarity.

³The file format should be easy to figure out. The file `words-d.txt` can be regarded as a matrix with about 70,000 rows and d columns. Each row is labeled by a word. The first line of the file is special: it gives the number of rows and columns.

⁴The details are not important for this assignment, but the vectors are optimized by gradient descent so as to arrange that the vector for each word token w_i will be predictable (to the extent possible) from the average vector of nearby word tokens (w_j for $j \neq i$ and $i - 5 \leq j \leq i + 5$). If you’re curious, you can find details in Mikolov et al. (2013)’s paper “Distributed Representations of Words and Phrases and their Compositionality,” available at <http://arxiv.org/abs/1301.3781>. We ran the CBOW method over the first 1 billion characters of English Wikipedia. `word2vec` doesn’t produce vector representations for rare words ($c(w) < 5$), so we first replaced rare words with the special symbol OOL (“out of lexicon”), forcing `word2vec` to learn a vector representation for OOL.

To measure the similarity of two words with vectors \vec{v} and \vec{w} , please use the *cosine similarity*, which is the cosine of the angle θ between the two vectors:

$$\cos \theta = \frac{\vec{v} \cdot \vec{w}}{\|\vec{v}\| \|\vec{w}\|} = \frac{\sum_{i=1}^d v_i w_i}{\sqrt{\sum_{i=1}^d v_i^2} \sqrt{\sum_{i=1}^d w_i^2}} \in [-1, 1]$$

What are the most similar words to **seattle**, **dog**, **communist**, **jpg**, **the**, and **google**? Play around some more. What are some examples that work “well” or “poorly”? What patterns do you notice?

So far you have been using $d = 50$. What happens for larger or smaller values of d ?

Hint: It’s probably a good idea to read the lexicon into some data structures.

Hint: Start by making **findsim** find only the single most similar word, which should be easy enough. Then you can generalize this method to find the 5 most similar words. (Since you only want the top 5, it’s not necessary to sort the whole lexicon by similarity—that method is acceptable, but inefficient.)

Note: You can copy the lexicon files to your personal machine. But to avoid downloading such large files, it may be easier to run **findsim** directly on the ugrad machines. If you do this, please don’t waste space by making fresh copies of the files on the ugrad machines. Just use them at their current location. If you like, create symbolic links (shortcuts) to them by typing “**ln -s /usr/local/data/cs465/hw-lm/lexicons/* .**” in your own working directory.

- (b) Now extend your program so that it can also be run as follows:

```
findsim words-50.txt king man woman
```

This should find the 5 words most similar to the vector **king** – **man** + **woman**, other than those three words themselves.

(The old command **findsim words-50.txt seattle** should still work. Note that it is equivalent to **findsim words-50.txt seattle seattle seattle**, and you may want to implement it that way.)

You can regard the above command as completing the analogy

man : woman :: king : ?


Try some more analogies, this time using the **200-dimensional vectors**. For example:

```
king - man + woman
paris - france + uk
hitler - germany + italy
child - goose + geese
goes - eats + ate
car - road + air
```

Come up with some analogy questions of your own. Which ones work well or poorly? What happens to the results if you use (say) $d = 10$ instead of $d = 200$?

Why does this work at all? Be sure to discuss the role of the vector **king** – **man** before **woman** is added. What does it tell you about the vectors that they can solve analogies?

9. **Extra credit:** Some politicians seem to speak in circles. Instead of uttering strings with a start and an end, they blow round linguistic objects out of their mouths, like smoke rings:



I wondered, how would we build a language model to describe the probability of such objects? Well, problems like this show up in computer vision, since images don't have a natural start and end either. Consider the similar problem of modeling the colors of these 4 pixels, using bigrams of adjacent pixel colors:

A	B
D	C

A well-known computer vision paper by J. Besag (1974) proposed approximating $p(A, B, C, D)$ by a product $p(A | B) \cdot p(B | C) \cdot p(C | D) \cdot p(D | A)$.

So here's the question. Can Besag's approximation be justified by using the chain rule plus backoff? If so, show how. If not, fix it as best you can. Discuss.

10. **Extra credit:** This problem is supposed to convince you that logical reasoning is just a special case of probabilistic reasoning.⁵ That is, probability theory allows you to draw a conclusion from some premises like “**If** there's no nail, **then** there's definitely no shoe.”

- | | | |
|-----|---|--|
| (a) | $p(\neg \text{shoe} \neg \text{nail}) = 1$ | <i>For want of a nail the shoe was lost,</i> |
| (b) | $p(\neg \text{horse} \neg \text{shoe}) = 1$ | <i>For want of a shoe the horse was lost,</i> |
| (c) | $p(\neg \text{race} \neg \text{horse}) = 1$ | <i>For want of a horse the race was lost,</i> |
| (d) | $p(\neg \text{fortune} \neg \text{race}) = 1$ | <i>For want of a race the fortune was lost,</i> |
| (e) | $p(\neg \text{fortune} \neg \text{nail}) = 1$ | <i>And all for the want of a horseshoe nail.</i> |

Show carefully that (e) follows from (a)–(d). *Hint:* Consider


$$p(\neg \text{fortune}, \neg \text{race}, \neg \text{horse}, \neg \text{shoe} | \neg \text{nail}),$$

as well as the “chain rule” and problems 1a, 1b, and 1k.

Note: The \neg symbol denotes the boolean operator NOT.

Note: Be glad I didn't ask you to prove the correct operation of Figure 1!

⁵An excellent and wide-ranging book developing this theme is *Probability Theory: The Logic of Science*, by the influential statistician E. T. Jaynes. See <http://bayes.wustl.edu/> for more readings.



Pencil Sharpener RUBE GOLDBERG (tm) RGI 038

Figure 1: RUBE GOLDBERG GETS HIS THINK-TANK WORKING AND EVOLVES THE SIMPLIFIED PENCIL-SHARPENER. Open window (A) and fly kite (B). String (C) lifts small door (D) allowing moths (E) to escape and eat red flannel shirt (F). As weight of shirt becomes less, shoe (G) steps on switch (H) which heats electric iron (I) and burns hole in pants (J). Smoke (K) enters hole in tree (L), smoking out opossum (M) which jumps into basket (N), pulling rope (O) and lifting cage (P), allowing woodpecker (Q) to chew wood from pencil (R), exposing lead. Emergency knife (S) is always handy in case opossum or the woodpecker gets sick and can't work.

The next discussion session will cover practice problems about probability and log-linear modeling, drawn from past exams. This is strongly recommended because those problems will help you apply these ideas to real situations. If you are unable to attend, you can ask me for the practice problems and solutions.